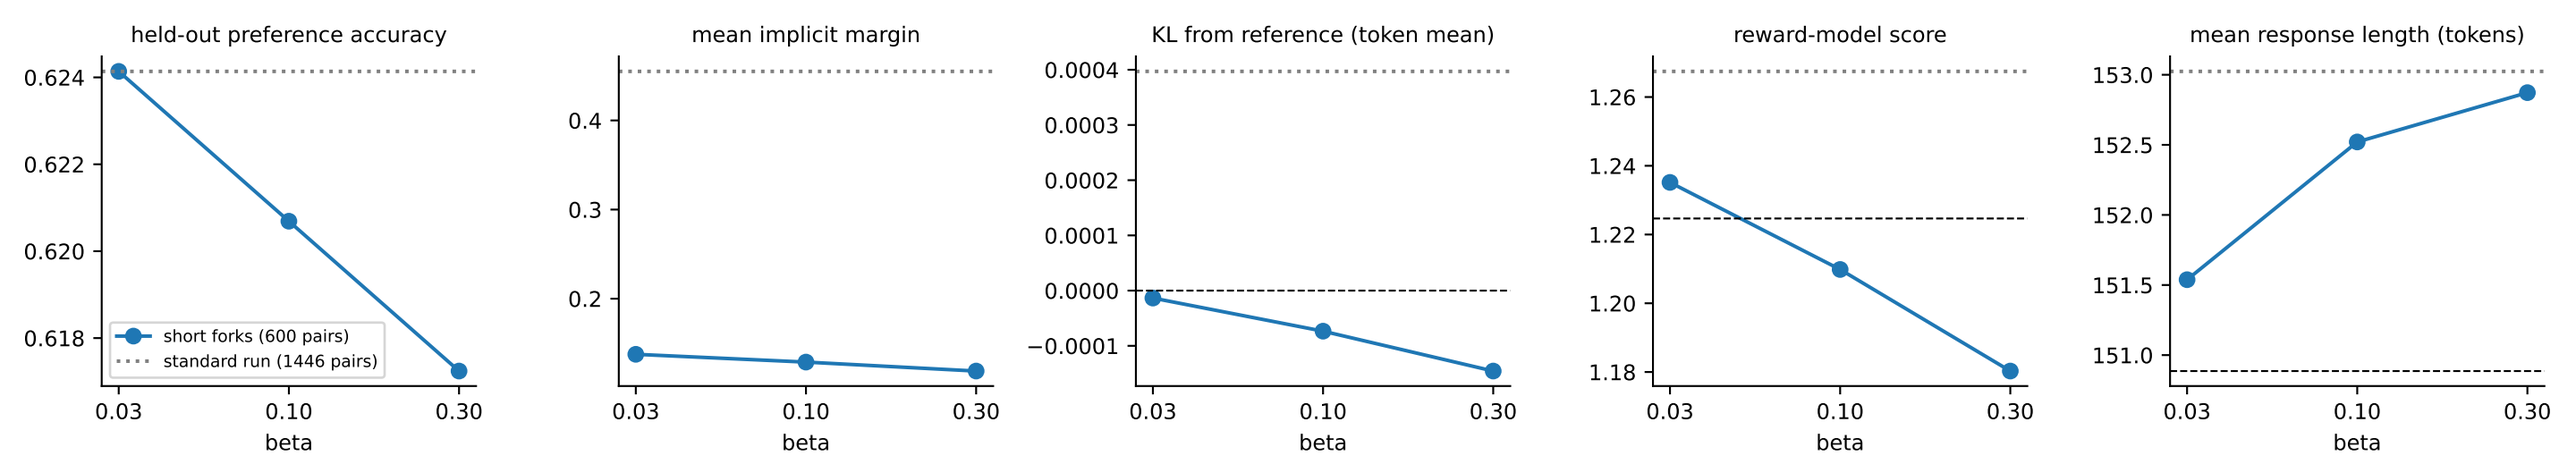
held-out preference accuracy
0.624
0.622
0.620
0.618
short forks (600 pairs)
standard run (1446 pairs)
0.03
0.10
0.30
beta
mean implicit margin
0.4
0.3
0.2
0.03
0.10
0.30
beta
KL from reference (token mean)
0.0004
0.0003
0.0002
0.0001
0.0000
−0.0001
0.03
0.10
0.30
beta
reward-model score
1.26
1.24
1.22
1.20
1.18
0.03
0.10
0.30
beta
mean response length (tokens)
153.0
152.5
152.0
151.5
151.0
0.03
0.10
0.30
beta