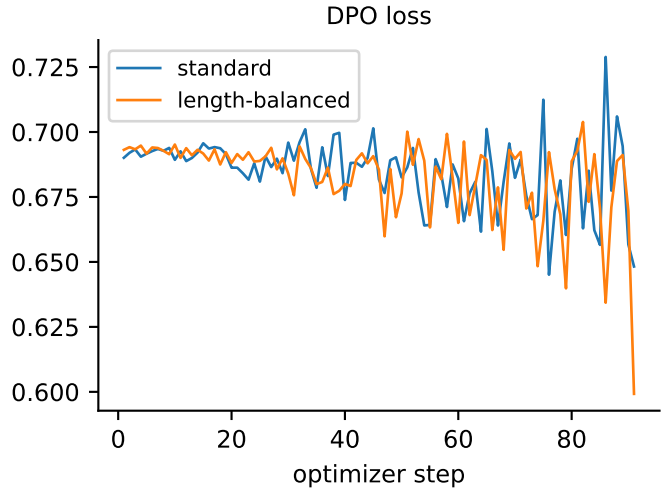
DPO loss
standard
length-balanced
0.725
0.700
0.675
0.650
0.625
0.600
0
20
40
60
80
optimizer step

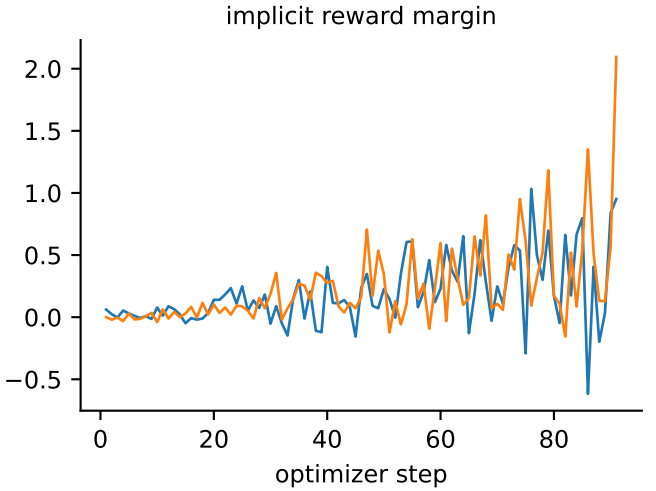
implicit reward margin
2.0
1.5
1.0
0.5
0.0
−0.5
0
20
40
60
80
optimizer step

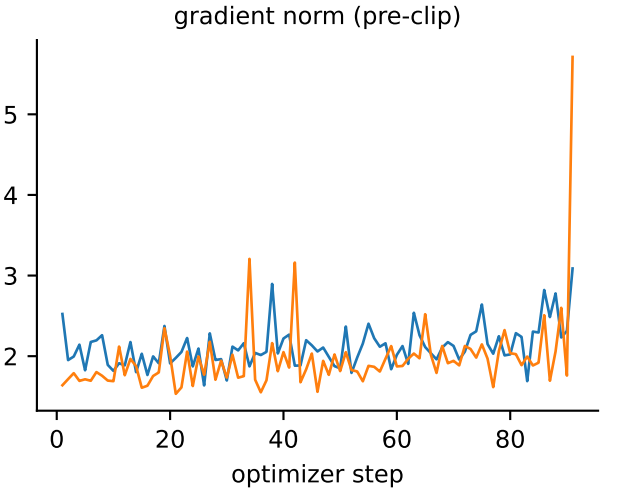
gradient norm (pre-clip)
5
4
3
2
0
20
40
60
80
optimizer step