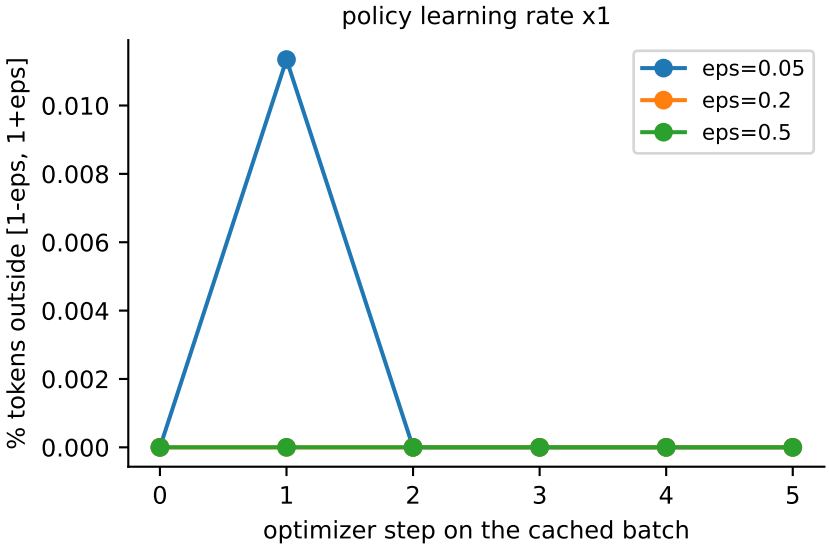
policy learning rate x1
% tokens outside [1-eps, 1+eps]
0.010
0.008
0.006
0.004
0.002
0.000
eps=0.05
eps=0.2
eps=0.5
0
1
2
3
4
5
optimizer step on the cached batch

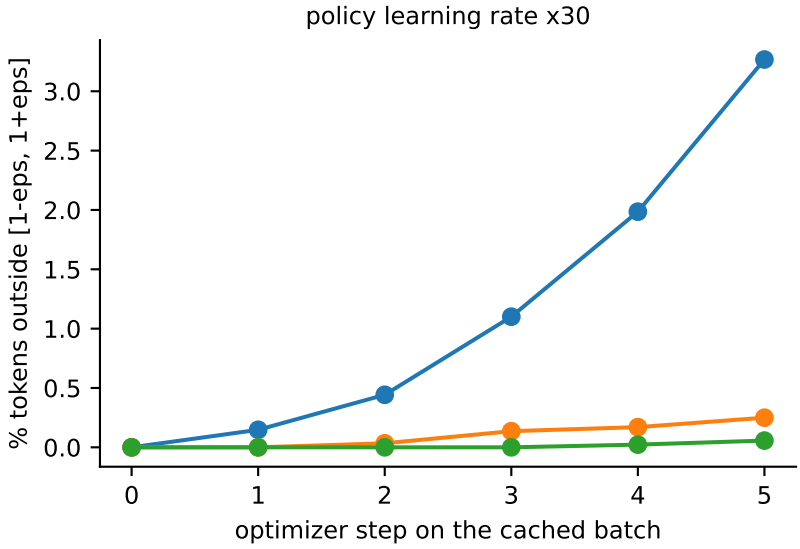
policy learning rate x30
% tokens outside [1-eps, 1+eps]
3.0
2.5
2.0
1.5
1.0
0.5
0.0
0
1
2
3
4
5
optimizer step on the cached batch